Gregory Joshua

Data Engineer | San Mateo, CA | gregjoshua11@gmail.com | (650) 924-8937 | LinkedIn | GitHub

EDUCATION

San Francisco State University | M.S. Business Analytics – GPA **3.83/4.00** **August 2024 – December 2026 (Expected)**

University of California, Santa Cruz | B.S. Mathematics – GPA **3.56/4.00** **August 2022 – August 2024**

EXPERIENCE

Kaiser Permanente *August 2026 - Present*

Graduate Analytics Consultant – SFSU Practicum *San Francisco, CA*

- Built a layered ELT pipeline with **DuckDB + dbt** to ingest and model historical data from **4 sources across 10+ years** into an analytics-ready warehouse, using SQL cases to address schema drift and dbt seeds for reference data.
- Implemented **100+ dbt tests** to validate nulls, uniqueness, referential integrity, required identifiers, and table grains; wrote documentation covering source-level data coverage gaps and their implications for downstream analysis.
- Explained **~70% of variance** in Kaiser's performance in the California individual marketplace by regressing market share on socioeconomic and market variables using **scikit-learn models**; presented findings to Kaiser stakeholders with **Tableau visualizations** in a PowerPoint deck.

Mathnasium *August 2024 - Present*

Data Analyst & Lead Instructor *San Francisco, CA*

- Wrote **Python ingestion scripts** to load student, session, and topic data from a proprietary instructional platform into **PostgreSQL**, creating a centralized warehouse for operational reporting and analysis.
- Built **dbt models** that transformed raw data into a **galaxy schema** and defined reusable metrics with **MetricFlow;** exposed the semantic layer using **Lightdash** replacing 10-20 minute manual Excel pulls with seconds of drag-and-drop analysis for center management.
- Orchestrated the end-to-end pipeline with **Apache Airflow**, automating Python ingestion, dbt transformations, and Lightdash deployment so newly generated instructional data flowed from the source system into refreshed reporting without manual intervention.

MedGenome *June 2021 – June 2022*

Data & Operations Analyst *Foster City, CA*

- Analyzed **LIMS and operational data** across **thousands of projects, samples, and client records** in Excel to answer recurring and ad-hoc questions from sales and operations; presented findings in weekly meetings with the **VP of Sales and operations stakeholders.**
- Combined project-tracking, LIMS, and sample-storage data to identify inefficiencies in the sample archiving process; redesigned and implemented the workflow onsite, **increasing storage efficiency by ~10%.**
- Coordinated the end-to-end flow of samples and associated project data across **clients, sales, research, and bioinformatics**, from sample receipt and storage through analysis and final data delivery, while maintaining project and compliance records.

PROJECTS

Retail Rocket ELT Pipeline and Data Warehouse **GitHub Repository**

- Built a **Dockerized ELT pipeline** that used **PySpark** to ingest and validate **23M+ e-commerce event and product-property records** into PostgreSQL, enforcing explicit schemas and source-data contracts before warehouse loading.
- Developed a warehouse using SQL window functions to sessionize **2.76M events into 1.76M sessions**, recursive CTEs to resolve product-category hierarchies, and latest-state logic for changing item properties.
- Implemented **pytest and dbt data-quality tests** across ingestion, staging, and marts, then defined governed business metrics with MetricFlow and exposed them through Lightdash for funnel, conversion, engagement, and cart-abandonment analysis.

TECHNICAL SKILLS

Data Engineering: Python (PySpark, pytest), SQL, PostgreSQL, DuckDB, dbt, MetricFlow, Apache Airflow, ETL/ELT, Docker, Git

Analytics: Lightdash, Tableau, Excel, KPI development, data visualization